



UNIVERSITÀ DI PISA

# Bright Computing Cluster Management

Presented By

Hassan Shabir and Ervisa Marku



# Mission Statement

- Bright Cluster Manager offers fast deployment and end-to-end management for heterogeneous high-performance computing (HPC) and AI server clusters at the edge, in the data center, and in multi/hybrid-cloud environments.
- It automates provisioning and administration for clusters ranging in size from a couple of nodes to hundreds of thousands, supports CPU-based and NVIDIA GPU-accelerated systems, and enables orchestration with Kubernetes.

# Bright Computing Inc

- Bright Computing, Inc. is a developer of software for deploying and managing high-performance (HPC) clusters, Kubernetes clusters, and OpenStack private clouds in on-premises data centers as well as in the public cloud.
- It was bought by NVIDIA in January 2022

# Building an HPC Cluster is Hard

- Most cluster management solutions use the “toolkit” approach (Linux distro + tools)
  - Tools typically used: Ganglia, Cacti, Nagios, xCAT, Chef, Puppet, Warewulf etc
- The problem with the “toolkit” approach :
  - Tools need to be glued together
  - Steep learning curve
  - Limited scalability
  - Tools use too many resources
  - Tools rarely designed for HPC
  - Hard to update the stack
- Making a collection of unrelated tools work together requires a lot of expertise and

# Quickly Build and Manage Heterogeneous Clusters



## **Automates the Building and Management of Clusters**

Heterogeneous high-performance Linux clusters can be quickly built and managed with NVIDIA Bright Cluster Manager, supporting HPC, machine learning, and analytics applications that span from core to edge to cloud.



## **Cross-Platform and Optimized for Accelerated Computing**

NVIDIA Bright Cluster Manager is ideal for heterogeneous environments, supporting Arm<sup>®</sup> and x86-based CPU nodes, and is fully optimized for accelerated computing with NVIDIA GPUs and NVIDIA DGX<sup>™</sup> systems.




## **Management Software that is Proven and Supported**

With thousands of deployments and enterprise support, NVIDIA Bright Cluster Manager is an ideal choice for deploying and managing HPC and AI clusters in production environments.

# About Bright Cluster Manager

- Bright Cluster Manager takes a much more fundamental & integrated approach
  - Single cluster management
  - Single database for configuration and monitoring data
  - Single CLI and GUI for all cluster management functionality
- Which makes Bright Cluster Manager
  - Easier to use, more scalable, secure, reliable, flexible and maintainable

# Cluster Structure



# Minimal Hardware Requirements

- **Head Node**

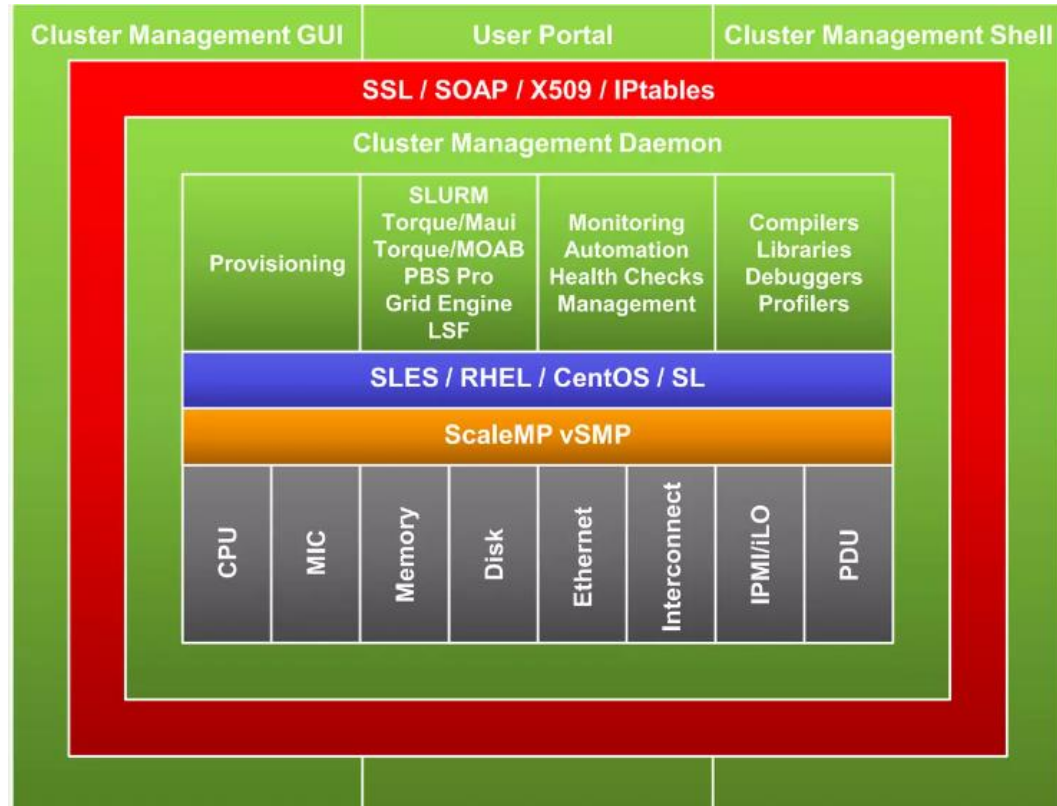
- x86-64 or Power8 CPU
- 4GB RAM
- 80GB diskspace
- 2 Gigabit Ethernet NICs (for the most common Type 1 topology (section 3.3.6))
- DVD drive or USB drive

- **Compute Nodes**

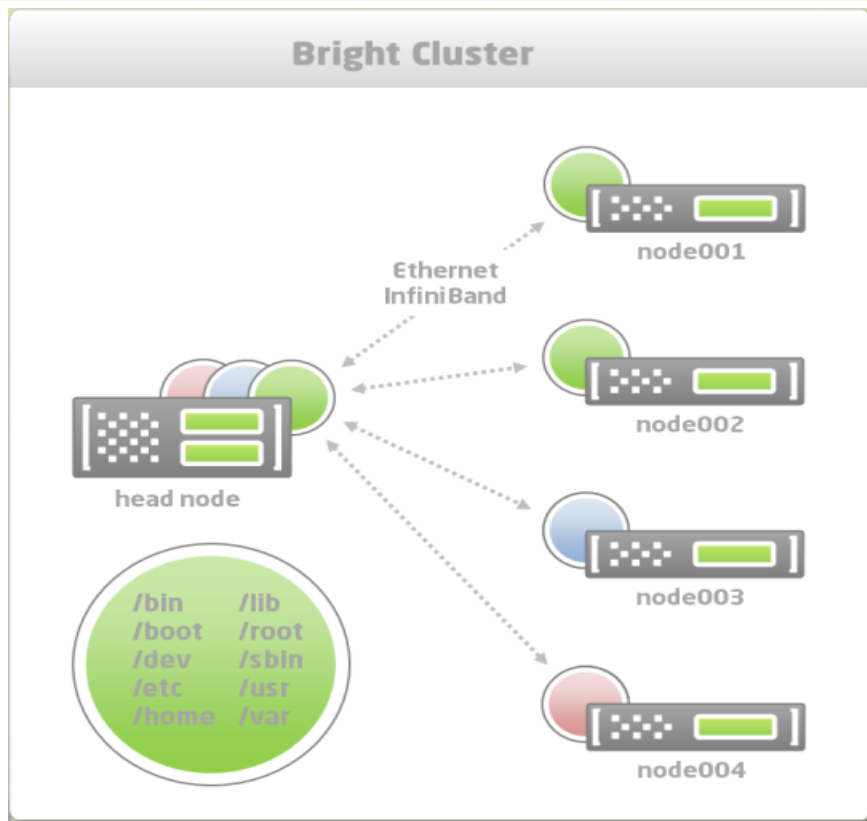
- x86-64 or Power8 CPU
- 1GB RAM (at least 4GB is recommended for diskless nodes)
- 1 Gigabit Ethernet NIC



# Bright Cluster Manager - Elements



# Provisioning



- Features

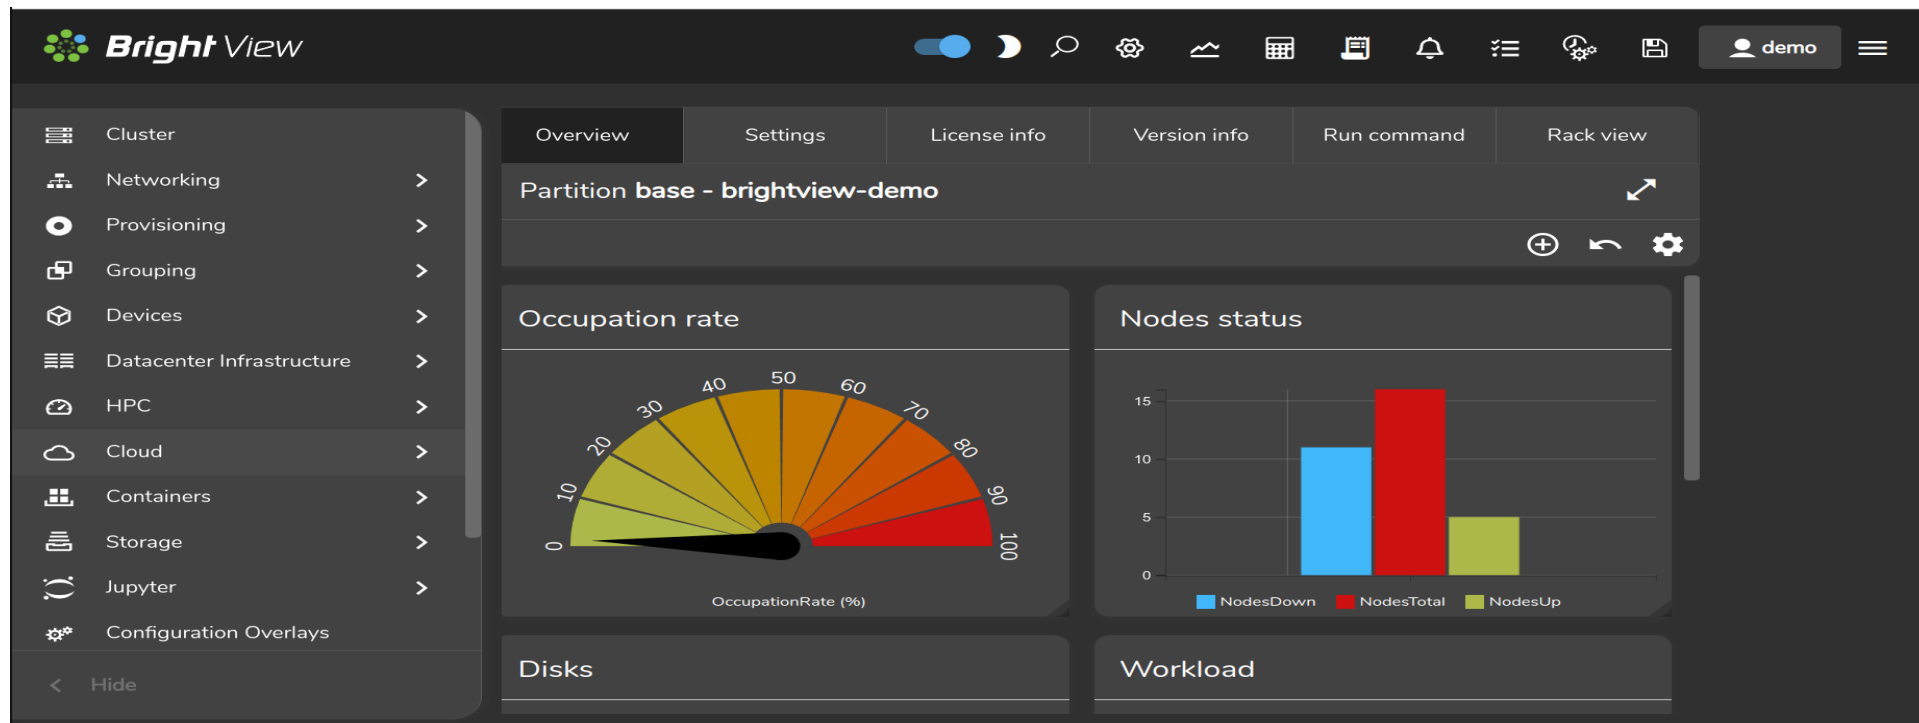
- Create: clone, from server
- Modify: `rpm -root`
- `Yum -installroot chroot`
- Sync: to and from
- Where: to HDD, diskless or VM
- How: update, full install, sync install
- Scalability: off-loadable, linear scaling

# Management Interface

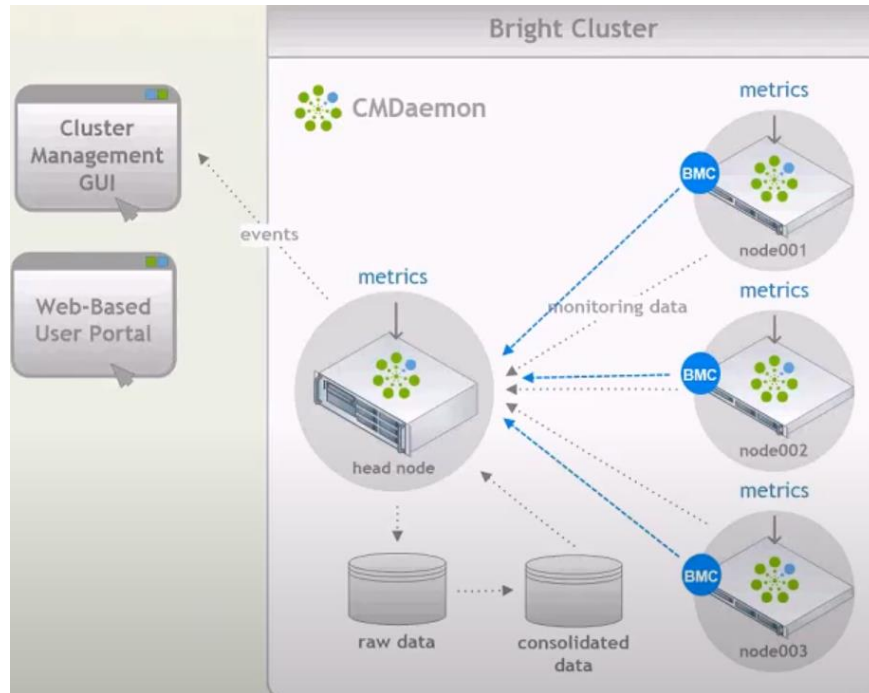
- **Graphical User Interface(CMGUI)**
  - Offers admins full cluster control
  - Standalone desktop application
  - Manages multiple clusters simultaneously
  - Runs natively on Linux, Windows and MacOSX
- **Cluster Management Shell(CMSH)**
  - All GUI functionality also available through Cluster Management Shell
  - Interactive and scriptable in batch mode



# GUI Demo

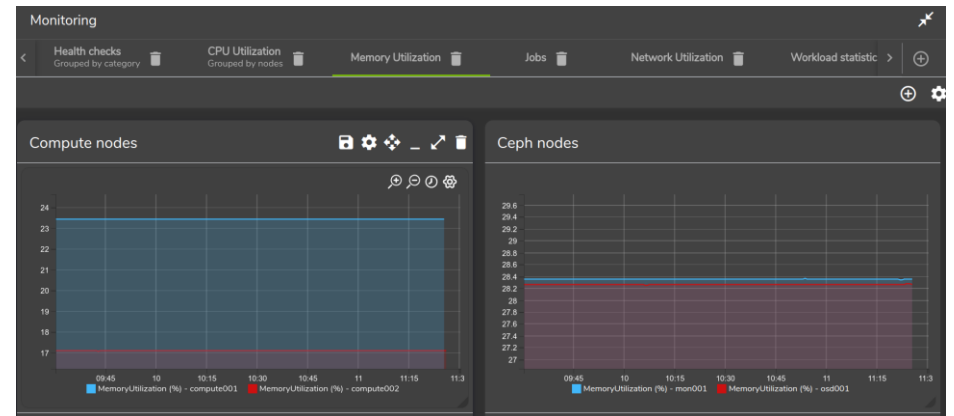
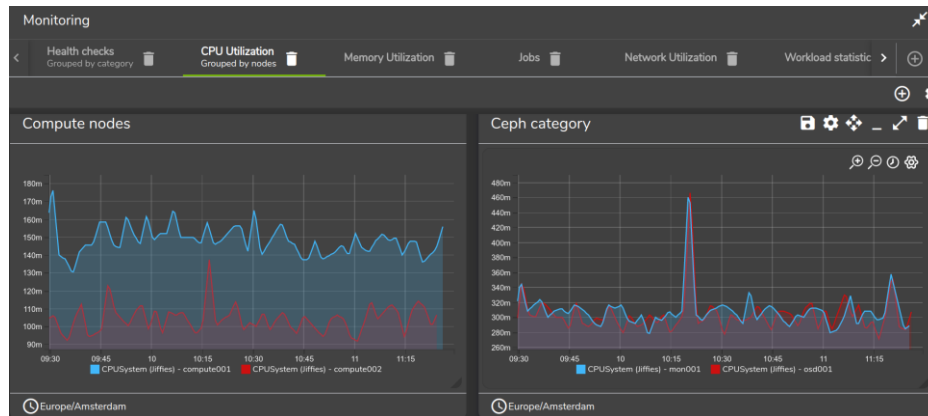


# Comprehensive Monitoring



- Monitor, visualize, and analyze a comprehensive set of hardware and software metrics with ease.

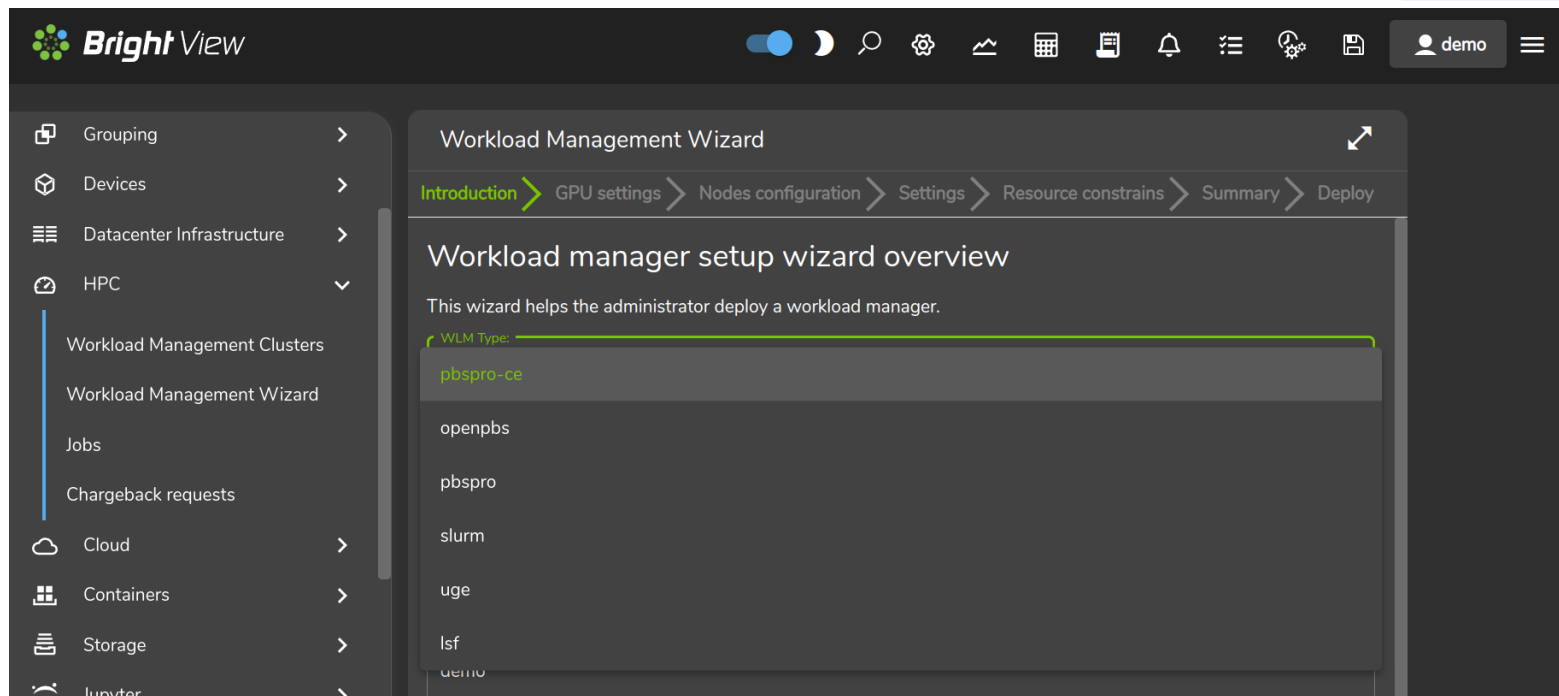
# Metrics



# Workload Manager Integration

- Automatic installation
- Automatic configuration
- Sampling, analysis and visualization of workload manager statistics
- Consistent GUI, User Portal and CLI front-end to workload manager
- Numeric GPU resource created
- Failover of workload manager
- Health checking
- Power saving through auto-power on/off based on workload

# HPC Workload Manager





## Storage Management – BeeGFS

- BeeGFS is a high-performance parallel file system, developed for HPC and optimized for intensive I/O.
- It uses a distributed metadata architecture designed for scalability and flexibility.
- Bright Cluster Manager provides packages to allow BeeGFS to be deployed, managed, and monitored on a Bright cluster.
- A BeeGFS cluster consists of a management server, one or more metadata servers and one or more storage servers.

# BeeGFS Integration

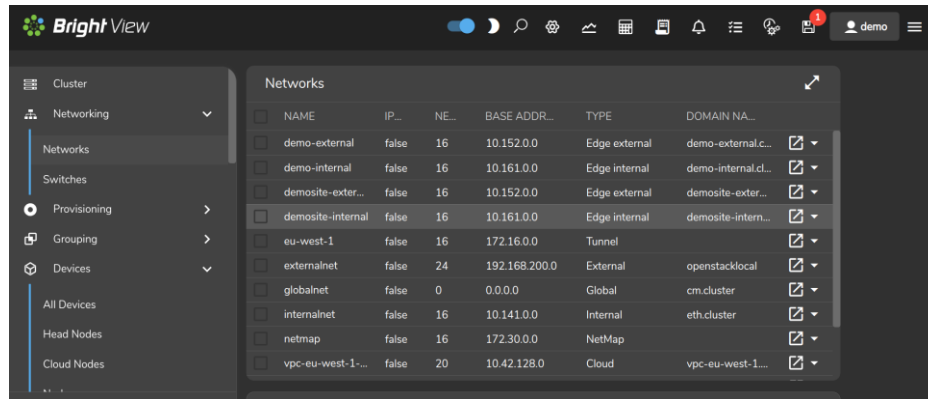
The screenshot displays the Bright View management console. On the left, a sidebar menu lists various system components: Containers (with sub-items Docker and Kubernetes), Storage (with sub-items Ceph and BeeGFS), and Jupyter. The 'BeeGFS' item is expanded, showing 'BeeGFS Wizard' and 'BeeGFS Configuration Overlays'. The main panel, titled 'Bee GFS Configuration Overlays', contains a table with the following data:

<input type="checkbox"/>	PRI...	NAME	NODES	CATEGORIES	ROLES	
<input type="checkbox"/>	500	beegfs-client	compute001,demo...		BeeGFS::Client	
<input type="checkbox"/>	500	beegfs-metadata	mon001		BeeGFS::Metadata	
<input type="checkbox"/>	500	beegfs-storage	osd001		BeeGFS::Storage	

At the bottom of the interface, there are buttons for 'ADD', 'REVERT', 'DELETE', and 'SAVE'.

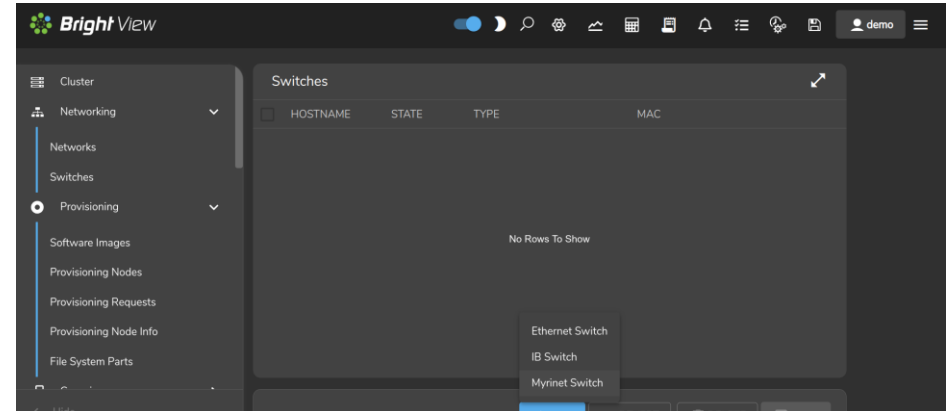
# Networking Management

- Support various network setups (VLAN, VxLAN, flat, etc.)



The screenshot shows the BrightView interface with the 'Networking' section selected in the sidebar. The 'Networks' tab is active, displaying a table of network configurations. The table has columns for NAME, IP, NE, BASE ADDR, TYPE, and DOMAIN NA. The data includes various network types such as Edge external, Edge internal, Tunnel, External, Global, Internal, NetMap, and Cloud.

NAME	IP	NE	BASE ADDR	TYPE	DOMAIN NA
demo-external	false	16	10.152.0.0	Edge external	demo-external.c...
demo-internal	false	16	10.161.0.0	Edge internal	demo-internal.c...
demosite-exter...	false	16	10.152.0.0	Edge external	demosite-exter...
demosite-internal	false	16	10.161.0.0	Edge internal	demosite-intern...
eu-west-1	false	16	172.16.0.0	Tunnel	
externalnet	false	24	192.168.200.0	External	openstacklocal
globalnet	false	0	0.0.0.0	Global	cm.cluster
internalnet	false	16	10.141.0.0	Internal	eth.cluster
netmap	false	16	172.30.0.0	NetMap	
vpc-eu-west-1...	false	20	10.42.128.0	Cloud	vpc-eu-west-1...



The screenshot shows the BrightView interface with the 'Networking' section selected in the sidebar. The 'Switches' tab is active, displaying a table of switch configurations. The table has columns for HOSTNAME, STATE, TYPE, and MAC. The data is currently empty, showing 'No Rows To Show'. A dropdown menu is visible at the bottom right, listing 'Ethernet Switch', 'IB Switch', and 'Myrinet Switch'.

HOSTNAME	STATE	TYPE	MAC
No Rows To Show			

# Comparisons with OpenHPC

- Functionality
- Ease of use
- Customization and flexibility
- Support and community
- Cost

•

# Conclusions

- **Bright Cluster Manager** provides a comprehensive, user-friendly solution for managing HPC clusters, covering both hardware and software aspects.
- **OpenHPC**, on the other hand, focuses on providing an open-source software stack for HPC environments, allowing for more customization but require more expertise.
- “The choice between the two depends on the specific needs, expertise, and resources available within an organization.”

